

Assignment is at the bottom!

```
In [1]: from sklearn.linear_model import LogisticRegression
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np


from pylab import rcParams
rcParams['figure.figsize'] = 20, 10

from sklearn.linear_model import LogisticRegression as Model
```

```
In [2]: y = np.concatenate([np.zeros(10), np.ones(10)])
x = np.linspace(0, 10, len(y))
```

```
In [3]: plt.scatter(x, y, c=y)
```

Out[3]: <matplotlib.collections.PathCollection at 0x23c6d905a50>



```
In [4]: model = LogisticRegression()
```

```
In [5]: model.fit(x.reshape(-1, 1), y)
```


Out[5]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [6]: plt.scatter(x,y, c=y)
plt.plot(x, model.predict_proba(x.reshape(-1, 1))[:,1])
```

Out[6]: [




```
In [7]: b, b0 = model.coef_, model.intercept_
model.coef_, model.intercept_
```

Out[7]: (array([[1.46709085]]), array([-7.33542562]))

```
In [8]: plt.plot(x, 1/(1+np.exp(-x)))
```

Out[8]: [




```
In [9]: b
```

```
Out[9]: array([[1.46709085]])
```

```
In [10]: plt.plot(x, 1/(1+np.exp(-(b[0]*x + b0))))
```

```
Out[10]: [ <matplotlib.lines.Line2D at 0x23c6fa00fa0>]
```



```
In [11]: from mpl_toolkits.mplot3d import Axes3D # noqa: F401 unused import


import matplotlib.pyplot as plt
from matplotlib import cm
from matplotlib.ticker import LinearLocator, FormatStrFormatter
import numpy as np

fig = plt.figure()
ax = fig.gca(projection='3d')

# Make data.
X = np.arange(-10, 10, 0.25)
Y = np.arange(-10, 10, 0.25)
X, Y = np.meshgrid(X, Y)
R = np.sqrt(X**2 + Y**2)
Z = 1/(1+np.exp(-(b[0]*X + b[0]*Y + b0)))
surf = ax.plot_surface(X, Y, Z, cmap=cm.coolwarm,
                      linewidth=0, antialiased=False)
```

C:\Users\gdlev\AppData\Local\Temp\ipykernel_14704\290434025.py:10: MatplotlibDeprecationWarning: Calling gca() with keyword arguments was deprecated in Matplotlib 3.4. Starting two minor releases later, gca() will take no keyword arguments. The gca() function should only be used to get the current axes, or if no axes exist, create new axes with default keyword arguments. To create a new axes with non-default arguments, use plt.axes() or plt.subplot().

```
ax = fig.gca(projection='3d')
```



In [12]: X

```
Out[12]: array([[ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                ...,
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75]])
```

In [13]: Y


```
Out[13]: array([[ -10.   ,  -10.   ,  -10.   , ...,  -10.   ,  -10.   ,  -10.   ],
                [ -9.75,  -9.75,  -9.75, ...,  -9.75,  -9.75,  -9.75],
                [ -9.5 ,  -9.5 ,  -9.5 , ...,  -9.5 ,  -9.5 ,  -9.5 ],
                ...,
                [  9.25,   9.25,   9.25, ...,   9.25,   9.25,   9.25],
                [  9.5 ,   9.5 ,   9.5 , ...,   9.5 ,   9.5 ,   9.5 ],
                [  9.75,   9.75,   9.75, ...,   9.75,   9.75,   9.75]])
```

What if the data doesn't really fit this pattern?

```
In [14]: y = np.concatenate([np.zeros(10), np.ones(10), np.zeros(10)])  
x = np.linspace(0, 10, len(y))
```

```
In [15]: plt.scatter(x,y, c=y)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x23c70676230>
```



```
In [16]: model.fit(x.reshape(-1, 1),y)
```


```
Out[16]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [17]: plt.scatter(x,y)
plt.plot(x, model.predict_proba(x.reshape(-1, 1)))
```

```
Out[17]: [<matplotlib.lines.Line2D at 0x23c706b1810>,
<matplotlib.lines.Line2D at 0x23c706b1870>]
```



```
In [18]: model1 = LogisticRegression()
model1.fit(x[:15].reshape(-1, 1),y[:15])
```

```
Out[18]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [19]: model2 = LogisticRegression()
model2.fit(x[15:].reshape(-1, 1), y[15:])
```


Out[19]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [20]: plt.scatter(x, y, c=y)
plt.plot(x, model1.predict_proba(x.reshape(-1, 1))[:, 1] * model2.predict_proba(x.
```

Out[20]: [<matplotlib.lines.Line2D at 0x23c70720f40>]



```
In [21]: df = pd.read_csv('../data/adult.data', index_col=False)
golden = pd.read_csv('../data/adult.test', index_col=False)
```

```
In [22]: from sklearn import preprocessing

enc = preprocessing.OrdinalEncoder()
```

```
In [23]: transform_columns = ['sex', 'workclass', 'education', 'marital-status',
                             'occupation', 'relationship', 'race', 'sex',
                             'native-country', 'salary']
```

```
In [24]: x = df.copy()

x[transform_columns] = enc.fit_transform(df[transform_columns])

golden['salary'] = golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K')
xt = golden.copy()

xt[transform_columns] = enc.transform(golden[transform_columns])
```



```
In [25]: df.salary.unique()
```

```
Out[25]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [26]: golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K').unique()
```

```
Out[26]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [27]: model.fit(preprocessing.scale(x.drop('salary', axis=1)), x.salary)
```

```
Out[27]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [28]: pred = model.predict(preprocessing.scale(x.drop('salary', axis=1)))
pred_test = model.predict(preprocessing.scale(xt.drop('salary', axis=1)))
```

```
In [29]: x.head()
```

```
Out[29]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	casualty
0	39	7.0	77516	9.0	13	4.0	1.0	1.0	4.0	1.0	
1	50	6.0	83311	9.0	13	2.0	4.0	0.0	4.0	1.0	
2	38	4.0	215646	11.0	9	0.0	6.0	1.0	4.0	1.0	
3	53	4.0	234721	1.0	7	2.0	6.0	0.0	2.0	1.0	
4	28	4.0	338409	9.0	13	2.0	10.0	5.0	2.0	0.0	

```
In [30]: from sklearn.metrics import (
    accuracy_score,
    classification_report,
    confusion_matrix, auc, roc_curve
)
```

```
In [31]: accuracy_score(x.salary, pred)
```

```
Out[31]: 0.8250360861152913
```

```
In [32]: confusion_matrix(x.salary, pred)
```

```
Out[32]: array([[23300, 1420],
               [ 4277, 3564]], dtype=int64)
```

```
In [33]: print(classification_report(x.salary, pred))
```

	precision	recall	f1-score	support
0.0	0.84	0.94	0.89	24720
1.0	0.72	0.45	0.56	7841
accuracy			0.83	32561
macro avg	0.78	0.70	0.72	32561
weighted avg	0.81	0.83	0.81	32561

```
In [34]: print(classification_report(xt.salary, pred_test))
```

	precision	recall	f1-score	support
0.0	0.85	0.94	0.89	12435
1.0	0.70	0.45	0.55	3846
accuracy			0.82	16281
macro avg	0.77	0.69	0.72	16281
weighted avg	0.81	0.82	0.81	16281

Assignment

1. Use your own dataset (create a train and a test set) and build 2 models: Logistic Regression and Decision Tree (shallow). Compare the test results using `classification_report` and `confusion_matrix`. Which algorithm is superior?

2. Repeat 1. but let the Decision Tree be much deeper to allow over-fitting. Compare the two models' test results again, and explain why it's superior

In [1]:

```
from sklearn import linear_model
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.tree import DecisionTreeClassifier
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import math
import re
plt.rcParams['figure.figsize'] = [10, 5]
from sklearn.metrics import (
    accuracy_score,
    f1_score,
    classification_report,
    confusion_matrix, auc, roc_curve,
    roc_auc_score
)
```

In [2]: `df_orig = pd.read_csv("NFL Play by Play 2009-2018 (v5).csv")`

```
C:\Users\gdlev\anaconda3\lib\site-packages\IPython\core\interactiveshell.py:333
1: DtypeWarning: Columns (42,166,167,168,169,174,175,178,179,182,183,188,189,19
0,191,194,195,203,204,205,218,219,220,231,232,233,238,240,241,249) have mixed t
ypes.Specify dtype option on import or set low_memory=False.
    exec(code_obj, self.user_global_ns, self.user_ns)
```

In [3]: `df = df_orig.copy()`

```
In [4]: all_cols = df.columns
cols_include = [
    "posteam_type",
    "yardline_100",
    "quarter_seconds_remaining",
    "half_seconds_remaining",
    "game_seconds_remaining",
    "game_half",
    "drive",
    # "sp",
    "qtr",
    "down",
    "goal_to_go",
    "ydstogo",
    "play_type",
    "shotgun",
    "no_huddle",
    "posteam_timeouts_remaining",
    "defteam_timeouts_remaining",
    "score_differential",
    "fourth_down_converted"
]
df = df_orig[cols_include].copy()
```

```

In [5]: # for i in df.columns:
#         print(f"Column: {i}")
#         print(df[i].value_counts())

# category columns
# posteam_type
# game_half
# play_type

### Predict fourth down conversions

df_4th = df[(df.down == 4) & df.play_type.isin(['pass', 'run'])].copy()
df_4th.loc[(df_4th.score_differential >= -3) & (df_4th.score_differential < 0),
df_4th.loc[(df_4th.score_differential == 0), 'curr_result'] = "tied"
df_4th.loc[(df_4th.score_differential <= -3) & (df_4th.score_differential >= -7),
df_4th.loc[(df_4th.score_differential > 0), 'curr_result'] = "winning"
df_4th.loc[(df_4th.score_differential <= -7), 'curr_result'] = "losing_td"
# df_4th = df_4th[(df_4th.score_differential >= -7) & (df_4th.score_differential
# # df_4th = df_4th[df_4th.half_seconds_remaining >= 500]
# # df_4th = df_4th[df_4th.game_seconds_remaining > 500]
# df_4th = df_4th[df_4th.goal_to_go == 0]
# df_4th = df_4th[df_4th.yardline_100 <= 50]
df_4th = df_4th[df_4th.ydstogo <= 2]

print(df_4th.describe())
print(df_4th.fourth_down_converted.value_counts())

cat_cols = ['posteam_type', 'game_half', 'curr_result']
df_4th = df_4th.drop(['down', 'play_type'], axis = 1)
df_4th = pd.concat([df_4th.drop(cat_cols, axis = 1), pd.get_dummies(df_4th[cat_cols])], axis = 1)
df_4th = df_4th.dropna()

# df_4th["curr_winning"] = np.where(df_4th.score_differential < 0, 0, 1)

```

	yardline_100	quarter_seconds_remaining	half_seconds_remaining	\
count	2321.000000	2321.000000	2321.000000	
mean	30.814735	385.418354	708.037915	
std	21.016028	260.931426	508.949574	
min	1.000000	1.000000	1.000000	
25%	12.000000	144.000000	231.000000	
50%	32.000000	360.000000	649.000000	
75%	45.000000	607.000000	1134.000000	
max	90.000000	900.000000	1800.000000	

	game_seconds_remaining	drive	qtr	down	goal_to_go	\
count	2321.000000	2321.000000	2321.000000	2321.0	2321.000000	
mean	1442.462732	13.308488	2.832831	4.0	0.139164	
std	1074.512766	7.344336	1.147550	0.0	0.346192	
min	1.000000	1.000000	1.000000	4.0	0.000000	
25%	418.000000	7.000000	2.000000	4.0	0.000000	
50%	1290.000000	14.000000	3.000000	4.0	0.000000	

75%	2384.000000	19.000000	4.000000	4.0	0.000000
max	3519.000000	34.000000	5.000000	4.0	1.000000

	ydstogo	shotgun	no_huddle	posteam_timeouts_remaining \
count	2321.000000	2321.000000	2321.000000	2321.000000
mean	1.233089	0.316674	0.059026	2.290392
std	0.422890	0.465279	0.235725	0.899779
min	1.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	2.000000
50%	1.000000	0.000000	0.000000	3.000000
75%	1.000000	1.000000	0.000000	3.000000
max	2.000000	1.000000	1.000000	3.000000

	defteam_timeouts_remaining	score_differential	fourth_down_converted
count	2321.000000	2321.000000	2321.000000
mean	2.435157	-4.306333	0.620422
std	0.836787	12.172666	0.485386
min	0.000000	-59.000000	0.000000
25%	2.000000	-12.000000	0.000000
50%	3.000000	-4.000000	1.000000
75%	3.000000	1.000000	1.000000
max	3.000000	48.000000	1.000000
1.0	1440		
0.0	881		

Name: fourth_down_converted, dtype: int64

```
In [6]: X_df = df_4th.drop("fourth_down_converted", axis = 1)
        y_df = df_4th[["fourth_down_converted"]]
```

```
In [7]: def get_accuracies(y_true, preds, model_name = "Model", verbose = True):
        acc = accuracy_score(y_true, preds)
        f1 = f1_score(y_true, preds)
        auc = roc_auc_score(y_true, preds)
        confus = confusion_matrix(y_true, preds)
        classif = classification_report(y_true, preds)

        acc_list = [acc, f1, auc, confus, classif]
        if verbose:
            print(f"\033[1m{model_name}\033[0m: \n\nAccuracy: {acc_list[0] * 100:.2f}

        return acc_list
```

```
In [128]: def train_model(X_df, y_df, trainsize = 0.75, model_type = "Logistic Regression",  
  
           X_train, X_test, y_train, y_test = train_test_split(X_df, y_df, train_size =  
           if model_type == "Logistic Regression":  
               model = linear_model.LogisticRegression()  
           else:  
               model = DecisionTreeClassifier(max_depth = maxdepth)  
  
           model.fit(X_train, y_train.values.ravel())  
           preds = model.predict(X_test)  
  
           get accuracies(y_test, preds, model_name = model_type)  
  
           return model
```

```
In [139]: print(train_model(X_df, y_df, model_type = "Logistic Regression"))
print(train_model(X_df, y_df, model_type = "Decision Tree"))
print(train_model(X_df, y_df, model_type = "Decision Tree", maxdepth = 10))
```

Logistic Regression:

Accuracy: 62.13%

F1-Score: 74.71%

AUC Score: 53.74%

Confusion Matrix:

```
[[ 36 190]
 [ 30 325]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.55	0.16	0.25	226
1.0	0.63	0.92	0.75	355
accuracy			0.62	581
macro avg	0.59	0.54	0.50	581
weighted avg	0.60	0.62	0.55	581

LogisticRegression()

Decision Tree:

Accuracy: 58.86%

F1-Score: 73.36%

AUC Score: 51.75%

Confusion Matrix:

```
[[ 13 233]
 [  6 329]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.68	0.05	0.10	246
1.0	0.59	0.98	0.73	335
accuracy			0.59	581
macro avg	0.63	0.52	0.42	581
weighted avg	0.63	0.59	0.46	581

DecisionTreeClassifier(max_depth=3)

Decision Tree:

Accuracy: 60.24%

F1-Score: 70.87%

AUC Score: 54.53%

Confusion Matrix:

```
[[ 69 149]
 [ 82 281]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.46	0.32	0.37	218
1.0	0.65	0.77	0.71	363
accuracy			0.60	581
macro avg	0.56	0.55	0.54	581
weighted avg	0.58	0.60	0.58	581

DecisionTreeClassifier(max_depth=10)

```
c:\Users\gdlev\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model\_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
In [118]: importance_df = pd.DataFrame(list(zip(model.feature_names_in_, model.feature_importances_)),
importance_df = importance_df.rename(columns = {importance_df.columns[0]: "feature_importance"})
```

```
In [66]: importance_df[importance_df.importance == 0].featurename
```

```
Out[66]: 5          qtr
6      goal_to_go
8      shotgun
9      no_huddle
10     posteam_timeouts_remaining
13      posteam_type_away
14      posteam_type_home
15      game_half_Half1
16      game_half_Half2
17      game_half_Overtime
18      curr_result_losing_fg
19      curr_result_losing_fg_td
20      curr_result_losing_td
21      curr_result_tied
22      curr_result_winning
Name: featurename, dtype: object
```

Compare the test results

According to this split, it appears that the logistic regression model slightly outperforms the shallow decision tree model. This is known because according to the classification report, the logistic regression model catches 92% of the cases where the possession team correctly converts a fourth down. In addition, the model predicts 63% correctly whether a team will or will not convert the fourth down. This leads to an overall F1 Score of 75%, F1 is a useful metric in comparing models, because it shows how prevalent false positives and false negatives that are predicted by the model. The model makes 581 predictions, the logistic regression model predicts a team should go for a fourth down 355 times, and predicts a team should kick or punt on the fourth down 226 times. In nominal terms, out of the 355 predictions to go for the fourth down, 325/355 are correct. This can be compared to the shallow decision tree model which has an overall F1 Score of 72.45%. Something to note however is that a false positive is significantly worse than a false negative. A false positive would represent the model predicting a team should go for a fourth down, even though they end up not converting, a false negative would be advising the team to kick or punt on the fourth down when they likely could have converted. The reason why a false positive is worse is because in football if a fourth down play is not converted than the opposing team immediately gets the ball where the play ends, this could lead to a significant disadvantage for the possession team. Due to this, the best model in my opinion, would be the model with the lowest amount of false positives which is decision tree which records $6/581 = 1\%$ which is less than the logistic regression model which has a false positive rate of $30/581 = 5\%$. ¶

Using the same logic as before, I would argue that the shallow decision tree model is even better than the deep decision tree model. Regardless of accuracy metrics, the most important aspect of the model is low false positive rate and the deep decision tree model actually has the highest false positive rate of any of the three models at $82/581 = 14\%$. However, I believe that if I created a more accurate model, that by increasing the depth of the decision tree would overfit the model and lead to much higher metrics, unfortunately I did not observe this phenomenon in my own data. Comparing all three models and using understanding of football strategy, I believe the model I would choose to use for this split would be the shallow decision tree model due to its low false positive rate.

